

# C Debugging & Optimization

A comprehensive guide on various aspects of software development, particularly focusing on **debugging**, **optimization**, and the use of certain GCC (GNU Compiler Collection) features to enhance both the performance and security of C programs.

## Source Code Optimization

### Use of `__builtin_unreachable()` in GCC

In GCC, there's a built-in function called `__builtin_unreachable()`. Your code should never call this function. The purpose of this is to mark a certain code block in your code as something that will never be executed, so the compiler can optimize it away.

**ASIDE:** Upcoming version **C23** has the function `unreachable()` which means the exact same, and this can be used in any compiler compliant with the spec.

```
int sqrt(int n) {
    // Suppose you know by precondition n is non-negative
    if (n < 0)
        __builtin_unreachable();
    /* code */
}
```

This is a message from you, the developer, to the compiler telling it to never worry about the case of `n` being negative. Of course, if you don't ensure the precondition or abuse `__builtin_unreachable()`, you could cause undefined behavior. **The compiler trusts that you know what you're doing.**

Notice how there's no need for conditional branching that could slow down the code. You use an `if` statement, but the compiler optimizes it away because of what it means. No code gets generated for that block.

You could also transform this into a macro:

```
#define assume(x) (x ? 0 : __builtin_unreachable())
```

p.s. A **macro** is a fragment of code which has been given a name. Whenever the name is used, it's replaced by the contents of the macro.

For example:

```
#define PI 3.14159

#define MAX(a, b) ((a) > (b) ? (a) : (b))
```

## Optimizing for cache(缓存) efficiency

A technique for arranging data and functions in memory to enhance performance by making better use of the CPU cache.

### 1. Data Alignment

It can fit variables into *cache lines* to maximize computations per cycle:

```
// This means allocate x at a multiple of 16
// Arg to aligned() must be a power of 2
long x __attribute__((aligned(16)));
```

`aligned(16)` ensures that the variable `x` is aligned in memory on a 16-byte boundary. This alignment can help reduce cache line misses.

A cache line miss occurs when the CPU cache must fetch data from the main memory (DRAM) because the required data is not in the cache. By aligning data to cache line boundaries, it will reduce the need for additional memory fetches and improving performance.

### 2. Function Attributes

```
// This function likely will NOT be called
int f(int) __attribute__((cold));
// The function likely WILL be called
int g(int) __attribute__((hot));
```

**hot** functions (`attribute((hot))`): By marking a function as "hot", you're indicating that it is likely to be executed frequently. The compiler can use this information to optimize the placement of the function's code in memory, keeping it in the instruction cache to ensure it's quickly accessible, minimizing instruction fetch times.

**cold** functions (`attribute((cold))`): Conversely, marking a function as "cold" suggests it's expected to be called infrequently. This allows the compiler to place this less critical code in a way that reduces its impact on the cache's efficiency for more frequently used data and code, potentially moving it out of the critical path in terms of cache and memory access.

The main point is that the code in the hot region will be *cached*. Cold code would sit in DRAM. This makes efficient use of our **instruction cache**.

## GCC Compilation Options for Security and Performance

### Security Improvements

**Stack overflow** is a programming error that occurs when a program uses more stack memory than is allocated.

We can simply protect against **stack overflow attacks**:

```
$ gcc -fstack-protector
```

`-fstack-protector` adds a **canary** -a randomly generated value- before the return address on the stack. This value is checked before returning from a function to ensure it hasn't been overwritten, which would indicate a stack overflow attack.

### Simpler Imagination of using `-fstack-protector`

You start by setting aside a random value (the canary) before doing anything risky. Then, after you're done, you check if that value got messed up by someone trying to attack your program. If it did, you know there's trouble, and you stop everything to prevent the attacker from causing harm.

This generates extra code in your function. It is as if you added something like:

```
int f() {
    // added check
    int i = randomvalue; // some macro probably

    char buf[512];
    read(fd, buf, 512);
    x = buf[12];

    // added check
    if (i != randomvalue)
        abort(); // reliably crash, dump core

    return x;
}
```

This means that attackers would have to guess the random value, meaning they'd have to guess from the possible values `int` can take on, and even more if something like `long` is used.

This works because the attacker would have to overwrite the `randomvalue` to get to the bottom of the call frame to reach the return address, which is typically what they need to overwrite to continue their attack.

There is debate whether `-fstack-protector` is a default option. Some Linux distros like Debian and Ubuntu have it on by default. This is off by default on SEASnet. There's also the opposite option called `-fno-stack-protector` that turns the setting off.

**ASIDE:** The **halting problem** states that there is no algorithm for looking at a program and figuring out if it would halt. There's no way to look at a program to figure out if it would stack overflow (Alan Turing).

Return Oriented Program (R.O.P.)

This is what the attacker trying to do. Which takes over return addresses and point them to your program somewhere else.

This tend to be slow, but useful to attackers. If they get enough return addresses, they can build a copied turing machine.

## Performance Improvements

### Performance optimization flags (-O, -Os, -Og, -O2, -O3, -O0)

Optimization flags are used by compilers to alter the way a program is compiled, which allow programmers to customize the tradeoffs among 1. debugging efficiency, 2. run time, and 3. compile time.

#### 1. -O0 (No optimization, default one)

**Debugging:** Easiest at this level because the code is compiled without any changes or optimizations. The program's structure closely mirrors the source code, so variable values and the flow of execution are predictable and match the source code, making it easier to step through with a debugger.

**Running Time:** Slowest, because no optimizations are applied that could speed up the execution. The program runs exactly as written, which might not be efficient.

**Compile Time:** Fastest, because the compiler does minimal work. It translates the source code directly into machine code without spending time analyzing the code for potential optimizations.

#### 2. -O1

**Debugging:** Still manageable, but slightly harder than at -O0 because the compiler has started to modify and optimize the code. Some optimizations may alter the flow of execution or optimize away variables, which can make debugging more challenging.

**Running Time:** Faster than -O0, as some basic optimizations are applied to make the program execute more efficiently.

**Compile Time:** Slightly slower than -O0 because the compiler does some analysis and optimization work, but generally not too much of a difference.

#### 3. -O2

**Debugging:** More difficult than -O1 because -O2 includes more aggressive optimizations. Variables may be optimized away, execution can be reordered, and some functions may be inlined, making it harder to trace through the program with a debugger.

**Running Time:** Generally much faster than -O1 because -O2 turns on a broad range of optimizations that focus on improving execution speed without regard to code size.

**Compile Time:** Longer than -O1 due to more complex optimizations being performed by the compiler, which requires more analysis and transformation of the code.

#### 4. -O3

**Debugging:** Most difficult at this level. The compiler applies all the optimizations of -O2 plus even more aggressive optimizations that can significantly change the structure of the code. This can make debugging with traditional step-through methods quite challenging.

**Running Time:** Potentially the fastest, although not all programs will see a performance gain from -O3 over -O2, and some may even run slower if the aggressive optimizations lead to poorer cache usage or other overheads.

**Compile Time:** Longest, as the compiler spends a lot of time performing extensive analysis to apply complex optimizations.

## 5. `-Os`

This flag enables optimizations that focus on reducing the size of the executable. The compile time can be longer than with `-O0` because the compiler is performing extra work to reduce the code size. The running time may be improved over `-O0`, especially on systems where memory bandwidth is a limiting factor, as a smaller executable can fit better in cache. However, debugging might be more challenging than with `-O0` because some optimizations may rearrange or eliminate code, making the debugging process less straightforward.

## 6. `-Og`

This flag is designed to provide a reasonable level of optimization without impeding the debugging process. The compile time will likely be longer than `-O0` because it includes more optimization passes, but typically shorter than `-Os` because it doesn't include all size reduction strategies. The running time will be shorter than with `-O0` due to the applied optimizations, but potentially longer than with `-Os` depending on the specific size-related optimizations that `-Os` might apply. Debugging is easier than with `-Os`, as `-Og` optimizations are chosen to preserve more debugging information.

The `-O`, `-O2`, `-O3`, etc. flags to specify the level of *\*optimization for CPU usage\**.

```
gcc -O foo.c
```

This is off by default because it slows down the compiler. The compiler spends more time finding ways to optimize the machine code to make the *runtime performance* faster. It's a trade off between compile time and runtime performance.

Additionally, the higher the level of optimization, the harder it becomes to debug the machine code because the compiler may choose to take certain liberties for the sake of making the executable more performant, making the machine code less parallel with the original source code.

By default, it cares about runtime performance, not executable size. To *optimize for size*, you use the `-Os` flag.

There is another optimization option that is still under construction `-Og`. This means to *optimize for debugging* i.e. make the code as "debuggable" as possible without affecting performance.

The `-O0` flag specifies to *not optimize at all*.

---

**ASIDE:** Consider this case where turning optimization on breaks the code:

```
double *p;  
*--p = *p++ * *p++;  
// "pop pop multiply push"
```

If you have competing side effects in the same statement, this is *undefined behavior*. The compiler is allowed to perform the actions in whatever order.

The compiler "activated" the bug by relying on undefined behavior in the name of optimization.

---

## Automating Hot/Cold with Profiling

1. You compile the program with a special flag that inserts extra code to count the number of times each instruction is executed. This understandably slows down execution.
2. You then run the program to gather the statistics.
3. Then, you recompile the program *with* the statistics. That way, the compiler knows which instructions are hot and cold and can decide how to mark them itself.

As long as your test runs are representative of how your program will actually be run, this is a great way of improving the performance of your application.

---

### Other GCC options

```
gcc --coverage
gprof
```

---

## Inlining

```
gcc -flto
```

Normally, GCC optimizes one module at a time.

```
// m.c
static int f(int x) {
    return -x;
}
```

Then suppose in the same module you call:

```
// m.c
f(3);
```

The compiler can deduce at compile-time that `f(3)` will always resolve to `-3`, so it can just expand it to `-3`. This is called **inlining**.

## Link-Time Optimization

**Link-time optimization (-flto)**: An advanced optimization technique that allows for optimizations across multiple compilation units, potentially significantly improving performance.

Then suppose you have definition and call in different modules:

```
// m.c
int f(int x) {
    return -x;
}
```

```
// n.c
extern int f(int x);
```

```
// n.c
f(3);
```

By default, the compiler can no longer make this expansion. It instead has to go through the overhead of a normal function call: putting `$3` into a register, etc.

The point of `gcc -flto` is to instruct the compiler to perform the optimization *across module boundaries* - whole program optimization.

```
gcc -flto foo.c bar.c baz.c
```

This creates a `foo.o`, `bar.o`, etc. that contains machine code, but within the file is also a copy of the source code. Thus, by the time you link them together:

```
gcc -flto foo.o bar.o baz.o
```

GCC has a copy of the entire source code, so it can perform the inline optimization to generate better code. The *downside* is that this command takes a long time. The optimization algorithms that GCC uses are around  $O(N^2)$ . It could take literal days with a large program.

## Catching Bugs with Static Checking

Debugging your program *before* the program runs. "Static" here has nothing to do with static allocation/static lifetime; it just means when the program isn't running. This is the opposite of **dynamic checking**.

- + No runtime overhead.
- + Covered bugs are 100% prevented by the time the program runs.
- - Some bugs cannot be statically checked for 100% reliably.

---

### 1. `assert()` macro

More usages in [Dynamic Checking](#).

`assert` is a macro used for debugging purposes, it tests a condition and, if the condition is false (evaluates to 0), it displays an error message on `stderr` and terminates the program by calling `abort()`.

For example:

```
#include <assert.h>
#include <stdio.h>

int main() {
    int a = 5;
    assert(a == 5); // This assertion passes, program continues.

    a = 3; // Change 'a' for demonstration purposes.
    assert(a == 5); // This assertion fails, error message is printed and
    program terminates.

    printf("This line is not executed.\n");
    return 0;
}
```

When `assert(a == 5);` is executed and `a` is not equal to 5, it will produce an error message like this:

```
Assertion failed: a == 5, file example.c, line 8
```

## 2. `static_assert()` macro

`static_assert()` is a macro introduced in C11 that performs **compile-time** assertion checks. Unlike the `assert()` macro, which evaluates conditions at **runtime**, `static_assert` checks conditions during compilation. If the condition fails, the compilation stops with a compile-time error.

```
#include <assert.h>

#define N 10 // Define N as a constant expression

int main() {
    static_assert(0 <= N, "N must be non-negative");
    return 0;
}
```

This use of `static_assert` ensures that `N` is non-negative at compile time. If it isn't, it's a compiler error.

**EXAMPLE:** Suppose you have a program that assumes a `long` is 64 bits:



```
#include <limits.h>
static_assert(LONG_MAX >> 31 >> 31 == 1);
```

This is once again like a *message from you to the compiler*. The compiler then checks beforehand for this condition before attempting to compile the program.

---

## GCC Warning Flags

In general, the flags beginning with **W** mean "warning". **-Wall** means "turn on all warnings":

```
$ gcc -Wall
```

However, people found that this tends to output extraneous/unnecessary warnings. Nowadays, **-Wall** has come to mean "turn on all warnings that most people will find useful."

**-Wall** implies:

- **-Wcomment**: warn about valid but bad comments like `/* a /* bad comment */`
- **-Wparentheses**: warn about style considered to be bad because something probably forgot parentheses in situations where operator precedence may be less familiar, like in `i << j + k` or `i < j || k < l && m < n`. The latter is common knowledge, but GCC disagrees, so this flag can be controversial.
- **-Waddress**: warn about making pointer comparisons when you probably didn't mean to, like `p = strchr("ab", "b"); if (p == "b") ...`
- **-Wstrict-aliasing**: warn about trying to "cheat" with pointers.

```
int i;
long *p = (long *)&i;
```

The C/C++ standard states that this results in undefined behavior, even if **int** and **long** happen to be the same size. This is controversial because a lot of low-level programs like the Linux kernel uses this all the time, *very carefully*. This is also why **casting** in general is risky because optimizing compilers may not do what you expect.

- **-Wmaybe-uninitialized**: warn about possible paths through a function where you use an uninitialized variable.

```
int f(int i)
{
    int n;
    if (i < 0)
```

```
    n = -i;
    return n; // if i >= 0, uninitialized n is returned
}
```

If you do something that checks out with arithmetic reasoning, the compiler may or may not be find with it, depending on other flags you pass it.

```
int f(int i)
{
    int n;
    if (i < 100)
        n = -1;
    if(i <= -1)
        return n;
}
```

Thus, there could be false positives if the compiler is not smart enough to deduce that the combination of things you've done always returns an initialized variable.

## GCC Extra Warning Flags

`-Wextra` is a collection of warnings like `-Wall` that are more controversial/less useful.

Some flags it includes are:

`-Wtype-limits`: warn about comparisons where the answer is obvious due to the type.

```
unsigned u;
if (u < 0)
    /* This code would never run! */
```

This is not always trivial. For example, if you're trying to run portable code, some `typedefs` might be different:

```
#include <time.h>
time_t t; // could be signed or unsigned
if (t < 0)
    /* This code may or may not run! */
```

In recent versions of GCC:

```
gcc -fanalyzer
```

This turns on **interprocedural warnings**. This looks at all callers and callees of a function to come up with a better picture of, say, if a variable is uninitialized or not. In other words, it looks through all *all paths through all calls* of every function (in current .c file, as that is what is deducible at compile-time). This is disabled by default because it slows the compiler down.

Specifying both of these flags would theoretically be rewarding but very expensive:

```
gcc -fanalyzer -flto
```

## Helping the Compiler through Source Code

You can modify your source code in minor ways to help the compiler do better checking.

### `_Noreturn`

```
// Prototype for some function that will never return
_Noreturn void fatal_error(char const *);

size_t size_sum(size_t a, size_t b)
{
    if (a <= SIZE_MAX - b)
        return a + b;
    fatal_error("size overflow");
}
```

Without the special `_Noreturn` signature, GCC will complain because it thinks `size_sum` is not always returning `size_t` even though it always does because `fatal_error` exits the program.

This will be standardized in C23 as `[[noreturn]]`.

`_Noreturn` also lets the compiler check whether a function declared to not return actually doesn't:

```
void fatal_error(char const *x)
{
    puts(x);
    exit(1);
}
```

`exit` is a `<stdlib.h>` function declared as something like:

```
_Noreturn void exit(int);
```

So `fatal_error` as it is defined above is fine. If we were to leave off `exit(1)`, then GCC would know that something is wrong.

## Pure Functions

With GCC, you can declare a **pure function** with the signature, telling the compiler that this function should not modify the state of the machine:

```
int hash(char *buf, size_t bufsize) __attribute__((pure));
```

This is stronger than merely declaring the parameters as **const**. A pointer to **const** just promises to not modify an object via its pointer, not that its a pointer to an actual constant.

This way, if you do something like:

```
char buf[100];
char c = buf[0];
i = hash(buf, sizeof buf);
// c == buf[0] must be true
// This gives the compiler opportunities to cache in register, etc.
```

This will be standardized in C23 as **[[reproducible]]**;

## Const Functions

Not to be confused with the **const** qualifier on C++ methods.

```
int square(int) __attribute__((const));
```

This is an even stronger flavor of pure functions. The return value can not even depend on the current state of the function, only its arguments.

A benefit is that if you call a **const** function multiple times with the same arguments, it could optimize knowing they would return the same value.

This will be standardized in C23 as **[[unsequenced]]**.

## Multiple Attributes

Example:

```
// ... is a syntax feature for variable number of arguments
int my_printf(char const *fmt, ...) __attribute__((nonnull(1) | format(printf, 1, 2)));
// nonnull(1) means the first argument cannot be NULL
// The format attribute does some type checking inside fmt arg:
// my_printf("a=%s", 1024); // not okay!
```

The names inside the `__attribute__` label, like `nonnull` and `format`, are not actual functions. They are *keywords* to the `__attribute__` syntax.

## Catching Bugs with Dynamic (Runtime) Checking

Quick review of `assert()`.

An example of dynamic checking, where you check at *runtime* and abort or similar if something fails:

```
#include <assert.h>
int f() {
    /* code */
    assert(0 <= n);
}
```

Checks you put in yourself. This is very flexible because it's your own code.

```
#include <stdckint.h> // C23 only
if (ckd_add(r, a, b))
    return EOVERFLOW;
// Not sure what this is but yeah
```

## GCC -fsanitize

Alternatively, you can let the compiler *insert* runtime checking.

```
gcc -fsanitize=undefined
```

Generate extra code to make the program reliably crash if it attempts undefined behavior (except for addresses). This causes the program to be slightly slower.

At the machine level, the adding overflow may be implemented like:

```
addl %eax,%ebx
jo   ouch      ; jump if overflow

ouch:
    call abort
```

This is the same as `-fsanitize=undefined`, but for addressing errors, such as subscripting an array. Due to technical reasons, this actually only catches *most* addressing errors, not all.

```
gcc -fsanitize=address
```

This attempts to insert runtime checks for **memory leaks**. Memory leaks are technically not undefined behavior nor errors.

```
gcc -fsanitize=leak
```

This attempts to insert runtime checks for **race conditions**, where there may be multiple threads that attempt to access the same I/O resource.

```
gcc -fsanitize=thread
```

Most of the time, these flags are useful for development only. They could be left in for production if the application prioritizes safety over performance.

## Valgrind

Valgrind is a programming tool for **memory debugging, memory leak detection, and profiling**. It can detect many memory-related errors that are common in C and C++ programs and that can lead to crashes and unpredictable behavior.

You can run a program called Valgrind on any program, no special compilation flags:  
And you are not running on the source code, but the program itself.

**Run Valgrind:** Use Valgrind to start your program. The simplest command is:

```
valgrind ./myprogram
```

This runs the program in a special environment and allows for more runtime checking.

**Analyze Output:** Valgrind will output information about memory leaks and errors it detects. For example, it might report memory that was allocated but never freed.

**Use Valgrind Tools:** Valgrind includes several tools; **memcheck** is the default and most widely used for memory error detection. You can specify a tool with the **--tool** option.

The summary screen displays information about memory leaks.

- **+** No special flags needed when compiling.
- **–** Much slower. Checking is also less extensive because Valgrind only has access to the machine code and not the underlying source code.

**BIG POINT:** Runtime checking is dicey. Even with all these tools, bugs can still slip through.

**ASIDE:** You can manually raise a compiler error with preprocessor directives in C:

```
#if INT_SIZE == LONG_SIZE
#error "ouch"
#endif
```

## Portability Checking

Portability is a large part of software construction issues.

### Portability Concerns

- **Architecture:** 32 vs. 64-bit platforms?

```
gcc          # default is 64-bit
gcc -m32     # check for 32-bit
```

- **OS:** Linux vs. MacOS vs. Androids vs. ...?
- **Software:** Chromium vs. Firefox vs. ...? (JS)
- **Software version:** Chromium v102 vs. v107 vs ...?

You need a good strategy to tackle this. One strategy is to have a comprehensive testing system. But you also need to consider the source code. The way it's done is defining a **portability layer**, a level of abstraction:

```
main code
  |
(some API)
  |
portability module
  |
Chromium, etc.
```

The compiler should be your servant, not your master. If you get warnings, look at what they're actually saying. If they're false positive, shut them off. - **Dr. Eggert, probably**

## Debugging Strategies

1. You're better off NOT debugging at all. Often times projects are stuck in **debugging hell**, spending more time debugging than actually developing. Debugging is an inefficient way of finding and fixing bugs. If you're spending a lot of time debugging, you should probably change your software construction approach such that you don't get so many bugs. *Be proactive* e.g. use static checking!
2. Write test cases. **Test-driven development (TDD)** is a theory of development that states that test cases are higher priority than code; one should write test cases first, the idea being you can use the test cases to debug the specifications of the code before actually writing the code.
3. Use a better platform.

4. **Defensive programming.** "When you're driving, assume that everyone is an idiot, drunk or both." Assume the other modules are broken.

1. You can use traces and logs (`print()` statements!).
2. Checkpoint/restart. Have one primitive in your program that saves the state of your entire application, probably to some file where it persists. Then have another primitive that restores the state:

```
save_state("foo") # do this periodically
restore_state("foo")
```

3. Assertions: crash when something that should never happen occurs.
4. Exception handling.
5. **Barricades**(路障).

```
+-----+-----+
| possibly | clean data |
| bad code | structs   |
+-----+-----+
      ^
  some barricade that processes
  code from the outside world, which
  you assume to be bad, before it
  makes it into your code
```

6. **Interpreters** (such as Valgrind) that execute code in a special environment.
7. **Virtual machines** that run a program in its own special sandbox, isolated from the outside system.

**TDD ASIDE:** If you write some test cases, and your program passes all the test cases, then you screwed up because you haven't found the bug you wanted to find. When you're writing test cases, you're trying to be imperfect. You're trying to think "how do I make this program crash?" Often times, tests are written by another class of developers because the self-interest of coders causes them to write bad test cases. - **Dr. Eggert, probably**

## Debugging Principles

- Don't guess! guessing doesn't scale to large program
- More systematic
  1. Stabilize the failure (reproduce the bug)
  2. Locate the failures cause/source
  3. Fix the bug

It is pretty common that the first two steps being most time-consuming



## Debugging Tools

Some examples are **GDB**, the GNU debugger, and **Valgrind**.

**GDB** is a portable debugger that works with many programming languages, including C and C++. It allows you to see what is going on inside another program while it executes or what another program was doing at the moment it crashed. GDB offers facilities for you to control the execution of the program, to inspect the values of variables, and to call functions independently of the program's normal behavior.

Key Features:

1. **Start and Stop Programs:** Control the execution of your program, specifying anything that might affect its behavior.
2. **Inspect Values:** Check the values of variables at specific points during program execution.
3. **Change Variables:** Modify variable values to test different execution paths.
4. **Call Functions:** Execute functions within the context of your running program, regardless of the program's normal behavior.
5. **Breakpoints and Watchpoints:** Stop the program execution at specified points or when variables change.

---

p.s. **strace** is a command that outputs the system calls a program uses, as if there were many **printf** logging calls in the source code.

```
#          vvvvvvvvvvvvvvvvvvvvv any shell command
strace cat /etc/os-release
```

Other commands we've seen before like **ps** and **top** are also *DevOps tools*. System administrators use this to monitor server activity, but they can also be useful for debugging.

---

## How Debuggers Work

GDB is actually a separate process from the process being debugged. It uses special system calls to exert some control over the process being debugged. Model:

```
(gdb process)---+
                  | special system calls
                  v
                (your process)
```

These special system calls include, at any point in execution:

- Starting/stopping/continuing

- Accessing memory
- Accessing registers

These do have security restrictions and cannot be used arbitrarily. These restrictions depend on the OS, but typically the rule is that the debugger must control a process with *same user ID*. Other OS may have a more restrictive rule, stating it can only debug a *child process*.

## Getting Started with Debugging: GCC

You can use a `-g` flag to specify debug info level for a C program:

```
gcc -g3 program.c
```

This conflicts with optimization because code that is optimized by the compiler tends to become harder to understand.

```
gcc -g3 -O2 program.c
```

This typically results in **inlining**, where calls to functions can be optimized away by substituting their body into places where it was called. The produced machine code would then be functionally equivalent but have lost the information that a function was even called.

What the `-g` flag does is bloat the resulting object and executable files with debugging tables. This data isn't visible when the program runs normally, but debuggers will be able to access them.

**ASIDE:** `_FORTIFY_SOURCE` is a standard technique used by GCC to make stack overflows less likely to succeed. For technical reasons, this is incompatible with no optimization i.e. attempting `gcc -O0` will cause a compiler error.

## Starting GDB

Dropping a program into GDB:

```
# vvvv program to debug
gdb diff
```

Now `gdb` is running but `diff` is not

1. After compiling your program with debug symbols, you start GDB by specifying the program's executable name. Inside GDB, you can set the working directory (`set cwd`) and environment variables (`set env`) for the program. This is crucial for programs that depend on specific directory paths or environment settings.

Setting the working directory of the program when it starts up:

```
(gdb) set cwd /etc
```

Setting environment variables for the debugging session: (This would apply to the program but not **gdb**)

```
(gdb) set env TZ America/New_York
```

**2. Address space layout randomization** (ASLR) is a security feature that randomizes memory addresses used by a program, making it harder for malicious exploits to predict the location of specific processes. However, this can make debugging unpredictable. GDB allows you to disable this feature (**set disable\_randomization**) to make debugging sessions more consistent.

A defense technique against buffer overflow attacks is to have the program run at randomized locations in memory (CS 33). By default, Linux executes programs in an environment with randomized addresses for the stack, heap, C library, etc. and many even the **main()** function.

The downside of this program is that it will run differently every time. This means that if there's a bug that depends on stack addresses for example, then it may appear sometimes and not for others. This makes debugging harder, so **by default**, this option is already **on**:

```
(gdb) set disable_randomization on (or off)
```

## Running the Program

1. In GDB, you can start your program with specific arguments using **run**. If your program is already running, you can attach GDB to it with its process ID (**attach**) and later release it using (**detach**) without stopping the program. To examine the program's call stack and understand how it reached its current state, use **bt** (backtrace).

Actually running the program. The arguments you supply after **run** are in shell syntax and forwarded to the executable being debugged:

```
(gdb) run -u /etc/os-release - < /dev/null
```

*Alternatively*, you can make GDB be in charge of another program using the PID of running process, effectively suspending it.

```
(gdb) attach 986317
```

Releasing the program:

```
(gdb) detach
```

Backtrace the program, which gets where the program is (like calling what functions, having what pointers)

```
(gdb) bt
```

## Controlling the Program

Control commands like **continue**, **step**, **next**, and their variants allow you to navigate through the program execution. These commands let you execute the program line by line, step into functions, or continue execution until the next breakpoint.

**^C** stops the program. GDB takes control.

**\*(int \*)0=27** crashes the program and falls under GDB's control.

### 1. **continue** or **c**

Resumes the execution of the program until it hits the **next breakpoint**, **watchpoint**, or until the program terminates.

Continue running the code:

```
(gdb) # c  
(gdb) continue
```

### 2. **step** or **s** (moves INTO function calls)

Executes the **next line** of code in the program. If the line contains a function call, step enters the function and pauses **at the first line of that function**.

Single step through the source code. Similarly, single step through the machine code.

```
(gdb) # s  
(gdb) step  
(gdb) # si  
(gdb) stepi
```

Stepping can be tricky because there isn't always a sequential mapping of source code lines to machine code lines. Stepping through some machine code lines may make it look like the program is jumping back and forth between source code lines instead of running one-by-one in order.

A courser-grained variant of the **step** command. Advancing to the next line of source code at the current function call level i.e. a single step but without worrying about function calls, stepping "over" them. Similarly, it has a machine code version.

### 3. **next** or **n** (steps OVER function calls)

Executes the next line of code similar to **step**, but if the line contains a function call, **next** executes the entire function as a single step and pauses when the function returns. Mainly use in the recursive function.

All these commands control the **instruction pointer**, e.g. **%rip** on x86-64.

```
(gdb) # n
(gdb) next
(gdb) # ni
(gdb) nexti
```

Finish the current function and then stop:

```
(gdb) fin
```

## Breakpoints and Watchpoints

1. Breakpoints are one of the most powerful features of GDB. You can set a breakpoint at specific functions or line numbers (**break**), causing the program to pause execution when it reaches these points. This lets you examine the program's state in detail.

You can use **breakpoints** to stop the program at a certain instruction, typically a function name. Creating a breakpoint:

```
(gdb) # b
(gdb) break analyze
(gdb) break diff.c:19
```

Listing your current breakpoints and their numbers:

```
(gdb) info break
```

Deleting a breakpoint by number:

```
(gdb) del 19
```

2. Watchpoints (**watch**) are similar but pause the program when the value of a specified variable changes.

You can use **watchpoints** to tell the program to run until the specified variable **p** changes value:

```
(gdb) watch p
```

### How does GDB implement breakpoints?

GDB takes the process being debugged and modifies its machine code. It stomps on the machine code of the specified function/line/instruction by zapping the first byte with a special instruction that is guaranteed to cause the program to trap, allowing GDB to take control.

### How does GDB implement watchpoints?

Single step through the code, and after each instruction, see if **p** has changed. This can be really slow unless you have special hardware support for watchpoints. Many CPUs, including x86-64, have this support.

## Extending GDB

```
define pl
  print *(long *) $arg1
end

(gdb) pl x
```

The point here is that debugging tool should be extensive and programable

- one example is **emacs/src/.gdbinit**

**GDB** supports either python or lisp

## Remote Debugging

If two machines have different architecture, there can be an issue that we need to translate the machine code from one to another

For example, from **x86-64** to **arm64**

## Other GDB Commands

Printing a C expression (or register values):

```
(gdb) # p
(gdb) print expr
(gdb) print $rax
(gdb) print a[5]
```

```
(gdb) print cos(3.0)

(gdb) p/x n (hexidecimal)
(gdb) p a[100]@10
(gdb) p x=y (be careful, you are actually changing the status)
```

It does more than just allow you to look at data. It lets you run a subroutine like `cos` in the program, which can modify the data and/or call arbitrary code from other parts of the program.

Disassembling a function to get the assembly code:

```
(gdb) disas cos
```

You can set a **checkpoint** and then run the code from the checkpoint by its number:

```
(gdb) checkpoint
(gdb) restart 42
```

The inverse of `continue` is `reverse-continue` (`rc`). This means to start running the program backwards until it hits the most recent breakpoint that it passed. This tends to be *very expensive* because GDB has to set a bunch of checkpoints under the hood.

```
(gdb) # rc
(gdb) reverse-continue
```

For **cross-debugging**, you can specify what target you want to run in

```
(gdb) target
```

This makes GDB run on some virtual machine or something?

```
(gdb) up/down
```

## GDB Example 1

### Step 1: Writing the Program

Create a file named `factorial.c` with the following code:

```
#include <stdio.h>

int factorial(int n) {
    if (n == 0) {
        return 1;
    } else {
        return n * factorial(n); // Bug: should be n-1
    }
}

int main() {
    int number = 5;
    printf("Factorial of %d is %d\n", number, factorial(number));
    return 0;
}
```

Notice the intentional bug in the `factorial` function: it calls `factorial(n)` instead of `factorial(n - 1)`, leading to an infinite recursion.

## Step 2: Compiling the Program with Debug Information

Compile the program with the `-g` flag to include debugging information:

```
$ gcc -g3 -o factorial factorial.c
```

## Step 3: Starting GDB

Start GDB with the compiled program:

```
$ gdb ./factorial
```

## Step 4: Setting a Breakpoint

Set a breakpoint at the `factorial` function to observe its behavior:

```
(gdb) break factorial
```

## Step 5: Running the Program

Run the program within GDB. It will start executing and stop when it reaches the `factorial` function:

```
(gdb) run
```



## Step 6: Stepping Through the Code

Once GDB hits the breakpoint, you can step through the code line by line to observe the behavior. Use the `next` command to go to the next line without stepping into other functions:

```
(gdb) next
```

However, since we're interested in seeing the recursive calls, use the `step` command to `step` into the recursive calls:

```
(gdb) step
```

After a few steps, it will become apparent that the function never reaches the base case (`n == 0`) because the decrement (`n - 1`) is missing, causing infinite recursion.

## Step 7: Identifying the Bug

At this point, you'll likely realize the bug: the function calls itself with the same value of `n` each time, instead of decrementing `n`. To confirm, you can print the value of `n` at each step:

```
(gdb) print n
```

## Step 8: Fixing the Bug

Exit GDB and fix the bug in the `factorial.c` file:

```
return n * factorial(n - 1);
```

## Step 9: Recompile and Re-test

After fixing the bug, recompile the program and test it again to ensure it now works correctly:

```
$ gcc -g3 -o factorial factorial.c  
$ ./factorial
```

## GDB Example 2

### Step 1: The Program with a Bug

Create a file named `sort.c` with the following code:

```
#include <stdio.h>

void bubbleSort(int arr[], int n) {
    int i, j, tmp;
    for (i = 0; i < n-1; i++) {
        for (j = 0; j < n-i-1; j++) {
            if (arr[j] > arr[j+1]) {
                // Swap the elements
                tmp = arr[j];
                arr[j] = arr[j+1];
                arr[j+1] = tmp;
            }
        }
    }
}

int main() {
    int arr[] = {64, 34, 25, 12, 22, 11, 90};
    int n = sizeof(arr)/sizeof(arr[0]);
    bubbleSort(arr, n);
    printf("Sorted array: \n");
    for (int i = 0; i < n; i++)
        printf("%d ", arr[i]);
    printf("\n");
    return 0;
}
```

## Step 2: Compiling with Debug Information

Compile the program with the `-g` flag to include debugging information:

```
$ gcc -g -o sort sort.c
```

## Step 3: Starting GDB

Launch GDB with the compiled program:

```
$ gdb ./sort
```

## Step 4: Setting Breakpoints

Set a breakpoint at the `bubbleSort` function to start debugging from the sorting logic:

```
(gdb) break bubbleSort
```

## Step 5: Running the Program

Start the program execution within GDB:

```
(gdb) run
```

When GDB hits the breakpoint, you're ready to start debugging the `bubbleSort` function.

## Step 6: Inspecting Variables and Stepping Through the Code

Use the `step` command to execute the code line by line, and use the `print` command to inspect variables:

```
(gdb) next
(gdb) print i
(gdb) print j
(gdb) print arr[j]
```

This allows you to observe how the variables change as the program executes, helping to identify where things might be going wrong.

**Step 7: Using Watchpoints** If you suspect that a variable is being changed unexpectedly, you can set a watchpoint on it:

```
(gdb) watch arr[j]
```

This will pause execution whenever `arr[j]` is modified, letting you examine the context of the change.

## Step 8: Continuing and Stepping into Functions

To proceed with execution until the next breakpoint or watchpoint, use the `continue` command. If you want to `step` into a function call (if there were any in this example), you could use the `step` command instead of `next` to descend into the function.

## Step 9: Viewing the Call Stack

If the program was more complex and called multiple functions, you might lose track of how you got to the current point of execution. Use the `backtrace` (or `bt`) command to view the call stack:

```
(gdb) backtrace
```

## Step 10: Conditional Breakpoints

Imagine you want to break only when a certain condition is true, such as when `i` is a specific value. Set a conditional breakpoint like this:

```
(gdb) break bubbleSort if i == 5
```

## Step 11: Finishing Function Execution

If you're in the middle of a function and want to run until it returns, use the `finish` command. This is useful for skipping over the rest of a function once you've seen what you need.

```
(gdb) finish
```